

The CURAND library

Will Landau

Iowa State University

November 4, 2013

Outline

Host interface

Device interface

Rejection sampling on the GPU

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

- ▶ **CURAND**: a CUDA C library for quickly generating pseudorandom and quasi-random numbers.
- ▶ **Pseudorandom sequence**: a sequence of numbers, generated by a deterministic algorithm, that has most of the properties of a truly random sequence.
- ▶ **Quasi-random (low-discrepancy) sequence**: a sequence of n -dimensional points, generated by a deterministic sequence, that appear random and appear to fill a region of n -dimensional space evenly.

Host and device APIs

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

▶ Host API

- ▶ Include the header, `curand.h`, and link with the `-lcuband` flag at compilation.
- ▶ Calls to number generators happen on the host.
- ▶ With each call, a predetermined number of random draws is generated and then stored for alter use in a kernel call or a copy statement.
- ▶ Supports 3 pseudorandom generators and 4 quasi-random generators.

▶ Device API

- ▶ Include the header, `curand_kernel.h`, and link with the `-lcuband` flag at compilation.
- ▶ Calls to number generators happen within kernels and device functions.
- ▶ Random numbers are generated and immediately used in real time on an as-need basis.
- ▶ For CUDA version 4.2, supports fewer generator algorithms then the host API.

Outline

Host interface

Device interface

Rejection sampling on the GPU

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

Using the host API

1. Create a new generator with `curandCreateGenerator()`.
2. Set the generator options. For example, use `curandSetPseudoRandomGeneratorSeed()` to set the seed.
3. Allocate memory for the random numbers with `cudaMalloc()`.
4. Generate random numbers with one or more calls to `curandGenerate()` or another generation function.
5. Clean up the generator with `curandDestroyGenerator()`.
6. Clean up everything else with `free()` and `cudaFree()`.

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

Generator types for `curandCreateGenerator()`

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

- ▶ Pseudorandom number generators
 - ▶ `CURAND_RNG_PSEUDO_DEFAULT`: XORWOW for the version currently on impact1
 - ▶ `CURAND_RNG_PSEUDO_XORWOW`: XORWOW algorithm
 - ▶ `CURAND_RNG_PSEUDO_MRG32K3A`: combined multiple recursive family
 - ▶ `CURAND_RNG_PSEUDO_MTGP32`: Mersenne Twister family
- ▶ Quasi-random number generators
 - ▶ `CURAND_RNG_QUASI_DEFAULT`: currently Sobol, 32-bit sequences
 - ▶ `CURAND_RNG_QUASI_SOBOL32`: Sobol, 32-bit sequences
 - ▶ `CURAND_RNG_QUASI_SOBOL64`: Sobol, 64-bit sequences
 - ▶ `CURAND_RNG_QUASI_SCRAMBLED_SOBOL32`: scrambled Sobol, 32-bit sequences
 - ▶ `CURAND_RNG_QUASI_SCRAMBLED_SOBOL64`: scrambled Sobol, 64-bit sequences

Generator options

- ▶ **Seed:** a 64-bit integer that initializes the starting state of a pseudorandom number generator.
- ▶ **Offset:** a parameter used to skip ahead in the sequence. Set $\text{offset} = 100$ to return the 100th number in the sequence first. Not available for the Mersenne Twister.
- ▶ **Order:** a parameter specifying how results are ordered in global memory.

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

Generator functions

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

► Random bits:

```
1 curandStatus_t curandGenerate(curandGenerator_t generator,  
    unsigned int *outputPtr, size_t num)
```

► Uniform(0, 1):

```
1 curandStatus_t curandGenerateUniform(curandGenerator_t generator,  
    float *outputPtr, size_t num)  
2  
3 curandStatus_t curandGenerateUniformDouble(curandGenerator_t  
    generator, double *outputPtr, size_t num)
```

► Normal:

```
1 curandStatus_t curandGenerateNormal(curandGenerator_t generator,  
    float *outputPtr, size_t n, float mean, float stddev)  
2  
3 curandStatus_t curandGenerateNormalDouble(curandGenerator_t  
    generator, double *outputPtr, size_t n, double mean,  
    double stddev)
```

Generator functions

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

► Log-normal:

```
1 curandStatus_t curandGenerateLogNormal(curandGenerator_t  
   generator, float *outputPtr, size_t n, float mean, float  
   stddev)  
2  
3 curandStatus_t curandGenerateLogNormalDouble(curandGenerator_t  
   generator, double *outputPtr, size_t n, double mean,  
   double stddev)
```

Example: host_api.cu

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

```
1  /*
2  * This program uses the host CURAND API to generate 10 pseudorandom
   floats.
3  */
4
5  #include <stdio.h>
6  #include <stdlib.h>
7  #include <cuda.h>
8  #include <curand.h>
9
10 int main(int argc, char *argv[]){
11     size_t n = 10;
12     size_t i;
13     curandGenerator_t gen;
14     float *devData, *hostData;
15
16     /* Allocate n floats on host */
17     hostData = (float *) calloc(n, sizeof(float));
18
19     /* Allocate n floats on device */
20     cudaMalloc((void **) &devData, n*sizeof(float));
21
22     /* Create a Mersenne Twister pseudorandom number generator */
23     curandCreateGenerator(&gen, CURAND_RNG_PSEUDO_MTGP32);
24
25     /* Set seed */
26     curandSetPseudoRandomGeneratorSeed(gen, 1234ULL);
27
28     /* Generate n floats on device */
29     curandGenerateUniform(gen, devData, n);
```

Example: host_api.cu

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

```

31  /* Copy device memory to host */
32  cudaMemcpy(hostData, devData, n * sizeof(float),
    cudaMemcpyDeviceToHost);
33
34  /* Show result */
35  printf("Random Unif(0, 1) draws:\n");
36  for(i = 0; i < n; i++) {
37      printf("  %1.4f\n", hostData[i]);
38  }
39  printf("\n");
40
41  /* Cleanup */
42  curandDestroyGenerator(gen);
43  cudaFree(devData);
44  free(hostData);
45  }

```

```

1  > nvcc host_api.cu -lcublas -o host_api
2  > ./host_api
3  Random Unif(0, 1) draws:
4  0.5823
5  0.4636
6  0.6156
7  0.9964
8  0.1182
9  0.2672
10 0.9241
11 0.7161
12 0.2309
13 0.4075

```

Outline

Host interface

Device interface

Rejection sampling on the GPU

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

Using the Device API

1. Within a kernel, call `curand_init()` to initialize the “state” of the random number generator.
 2. Within a (possibly separate) kernel, call `curand()` or one of its wrapper functions (such as `curand_uniform()` or `curand_normal()`) to generate pseudorandom or quasi random numbers as needed.
- ▶ RNG types available
 - ▶ Pseudorandom
 - ▶ XORWOW
 - ▶ Quasi-random
 - ▶ 32-bit Sobol
 - ▶ 32-bit scrambled Sobol

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

Device API functions: XORWOW

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

```
1  __device__ void curand_init (unsigned long long seed ,
2                               unsigned long long sequence ,
3                               unsigned long long offset ,
4                               curandState_t *state)
5
6  __device__ unsigned int
7  curand (curandState_t *state) // RANDOM BITS
8
9  __device__ float
10 curand_uniform (curandState_t *state) // U(0,1)
11
12 __device__ double
13 curand_uniform_double (curandState_t *state) // U(0,1)
14
15 __device__ float
16 curand_normal (curandState_t *state) // N(0,1)
17
18 __device__ double
19 curand_normal_double (curandState_t *state) // N(0,1)
```

Device API functions: XORWOW

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

```
1  __device__ float2
2  curand_normal2 (curandState_t *state) // 2 N(0,1) draws
3
4  __device__ float2
5  curand_log_normal2 (curandState_t *state) // 2 N(0,1) draws
6
7  __device__ float
8  curand_log_normal (curandState_t *state, float mean, float stddev)
9
10 __device__ double
11 curand_log_normal_double (curandState_t *state, double mean, double
    stddev)
12
13 __device__ double2
14 curand_normal2_double (curandState_t *state) // 2 draws
15
16 __device__ double2
17 curand_log_normal2_double (curandState_t *state) // 2 draws
```


Device API functions: Sobol

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

```
1  __device__ void
2  curand_init (
3      unsigned int *direction_vectors ,
4      unsigned int offset ,
5      curandStateSobol32_t *state) // Sobol
6
7  __device__ void
8  curand_init (
9      unsigned int *direction_vectors ,
10     unsigned int scramble_c ,
11     unsigned int offset ,
12     curandStateScrambledSobol32_t *state) // Scrambled Sobol
13
14 __device__ unsigned int
15 curand (curandStateSobol32_t *state)
16
17 __device__ float
18 curand_uniform (curandStateSobol32_t *state)
```

Device API functions: Sobol

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

```
1  __device__ float
2  curand_normal (curandStateSobol32_t *state)
3
4  __device__ float
5  curand_log_normal (
6      curandStateSobol32_t *state ,
7      float mean,
8      float stddev)
9
10 __device__ double
11 curand_uniform_double (curandStateSobol32_t *state)
12
13 __device__ double
14 curand_normal_double (curandStateSobol32_t *state)
15
16 __device__ double
17 curand_log_normal_double (
18     curandStateSobol32_t *state ,
19     double mean,
20     double stddev)
```

Example: device_api.cu

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

```

1  /*
2  * This program uses the device CURAND API to calculate what
3  * proportion of pseudo-random ints are odd.
4  */
5
6  #include <stdio.h>
7  #include <stdlib.h>
8  #include <cuda.h>
9  #include <curand_kernel.h>
10
11 __global__ void setup_kernel(curandState *state){
12     int id = threadIdx.x + blockIdx.x * 64;
13
14     /* Each thread gets same seed, a different sequence number, no
15        offset */
16     curand_init(1234, id, 0, &state[id]);
17 }
18
19 __global__ void generate_kernel(curandState *state, int *result){
20     int id = threadIdx.x + blockIdx.x * 64; int count = 0;
21     unsigned int x;
22
23     /* Copy state to local memory for efficiency */
24     curandState localState = state[id];
25
26     /* Generate pseudo-random unsigned ints */
27     for(int n = 0; n < 100000; n++){
28         x = curand(&localState);

```

Example: device_api.cu

```
28     /* Check if odd */
29     if(x & 1){
30         count++;
31     }
32 }
33
34 /* Copy state back to global memory */
35 state[id] = localState;
36
37 /* Store results */
38 result[id] += count;
39 }
40
41 int main(int argc, char *argv[]){
42     int i, total;
43
44     int *devResults, *hostResults;
45     curandState *devStates;
46
47     /* Allocate space for results on host */
48     hostResults = (int *) calloc(64 * 64, sizeof(int));
49
50     /* Allocate space for results on device */
51     cudaMalloc((void **)&devResults, 64 * 64 * sizeof(int));
52
53     /* Set results to 0 */
54     cudaMemset(devResults, 0, 64 * 64 * sizeof(int));
55
56     /* Allocate space for prng states on device */
57     cudaMalloc((void **)&devStates, 64 * 64 * sizeof(curandState));
```

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

Example: device_api.cu

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

```

59  /* Setup prng states */
60  setup_kernel<<<64, 64>>>(devStates);
61
62  /* Generate and use pseudorandom numbers*/
63  for(i = 0; i < 10; i++){
64      generate_kernel<<<64, 64>>>(devStates, devResults);
65  }
66
67  /* Copy device memory to host */
68  cudaMemcpy(hostResults, devResults, 64 * 64 * sizeof(int),
             cudaMemcpyDeviceToHost);
69
70  /* Show result */
71  total = 0;
72  for(i = 0; i < 64 * 64; i++) {
73      total += hostResults[i];
74  }
75  printf("Fraction odd was %10.13f\n", (float) total / (64.0f * 64.0f
             * 100000.0f * 10.0f));
76
77  /* Cleanup */
78  cudaFree(devStates);
79  cudaFree(devResults);
80  free(hostResults);
81
82  return EXIT_SUCCESS;
83  }

```

Example: device_api.cu

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

```
1 > nvcc device_api.cu -lcublas -o device_api
2 ptxas /tmp/tmpxft_000020d0_00000000-2_device_api.ptx, line 501;
3 warning : Double is not supported. Demoting to float.
4 >
5 > ./device_api
6 Fraction odd was 0.4999966323376
```

Outline

Host interface

Device interface

Rejection sampling on the GPU

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

RNG types supported

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

- ▶ Dr. Jarad Niemi wrote example rejection sampling code available at <https://github.com/jarad/gpuRejectionSampling>.
- ▶ Idea
 1. Draw a pseudorandom number, x .
 2. If x is too big, throw out x and return to step 1.
 3. Return x .

cpu_runif.c

```

1 #include <Rmath.h>
2 // #include <stdlib.h>
3
4
5 int cpu_runif(int n, double ub, int ni, int nd, double *u, int *count)
6 {
7     int i, j, a;
8     double b;
9     GetRNGstate();
10    for (i=0; i<n; i++) {
11        count[i] = -1;
12        u[i] = ub+1;
13
14        while ( u[i]>ub ) {
15            count[i]++;
16            // u[i] = rand() / ((double)RAND_MAX + 1);
17            u[i] = runif(0,1);
18
19            // Computational overhead
20            a=0; for (j=0; j<ni; j++) a += 1;
21            b=1; for (j=0; j<nd; j++) b *= 1.00001;
22        }
23    }
24    PutRNGstate();
25 }
26
27 void cpu_runif_wrap(int *n, double *ub, int *ni, int *nd, double *u,
28                    int *count){
29     cpu_runif(*n, *ub, *ni, *nd, u, count);
30 }

```

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

gpu_runif.cu

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

```

1 #include <curand_kernel.h>
2 #include "cutil_inline.h"
3
4 #define THREADS_PER_BLOCK 256
5
6 __global__ void setup_prng(unsigned long long seed, curandState *
    state)
7 {
8     int id = threadIdx.x + blockIdx.x * THREADS_PER_BLOCK;
9     curand_init(seed, id, 0, &state[id]);
10 }
11
12 __global__ void runif_kernel(curandState *state, double ub, int ni,
    int nd,
13                             double *uniforms, int *counts)
14 {
15     int i, a, count, id = threadIdx.x + blockIdx.x *
        THREADS_PER_BLOCK;
16     double b, u;
17
18     // Copy state to local memory for efficiency */
19     curandState localState = state[id];
20
21     // Find random uniform below the upper bound
22     count = -1;
23     u = ub+1;

```

gpu_runif.cu

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

```

24 while ( u>ub )
25 {
26     count++;
27     u = curand_uniform_double(&localState);
28
29     // Computational overhead
30     a=0; for (i=0; i<ni; i++) a += 1;
31     b=1; for (i=0; i<nd; i++) b *= 1.00001;
32 }
33
34 // Copy state back to global memory */
35 state[id] = localState ;
36
37 // Store results */
38 uniforms[id] = u;
39 counts[id] = count;
40 }
41
42 //CURAND_RNG_PSEUDO_MTGP32
43
44 extern "C" {
45
46 void gpu_runif(int *n, double *ub, int *ni, int *nd, double *seed ,
47               double *u, int *c)
48 {
49     int nBlocks = *n/THREADS_PER_BLOCK, *d_c;
50     size_t u_size = *n * sizeof(double), c_size = *n * sizeof(int);
51     double *d_u;
52
53     cutiSafeCall( cudaMalloc((void**)&d_u, u_size) );
54     cutiSafeCall( cudaMalloc((void**)&d_c, c_size) );

```

gpu_runif.cu

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

```

54 // Setup prng states
55 curandState *d_states;
56 cutilSafeCall( cudaMalloc((void**)&d_states, nBlocks*
57     THREADS.PER_BLOCK*sizeof(curandState)) );
58 setup_prng<<<nBlocks,THREADS.PER_BLOCK>>>(*seed, d_states);
59 runif_kernel<<<nBlocks,THREADS.PER_BLOCK>>>(d_states, *ub, *ni, *
60     nd, d_u, d_c);
61 cutilSafeCall( cudaMemcpy(u, d_u, u_size,
62     cudaMemcpyDeviceToHost) );
63 cutilSafeCall( cudaMemcpy(c, d_c, c_size,
64     cudaMemcpyDeviceToHost) );
65 cutilSafeCall( cudaFree(d_u) );
66 cutilSafeCall( cudaFree(d_c) );
67 cutilSafeCall( cudaFree(d_states) );
68 }
69 } // end of extern "C"

```

my.unif.r

```

1 my.runif = function(n, ub, ni=1, nd=1,
2                     engine="R", seed=1)
3 {
4   engine = pmatch(engine, c("R", "C", "GPU"))
5
6   switch(engine,
7   {
8     # R implementation
9     u = rep(Inf, n)
10    count = rep(0, n)
11    set.seed(seed)
12
13    for (i in 1:n) while( (u[i] <- runif(1))>ub )
14    {
15      count[i] = count[i]+1
16      a = 0
17      b = 1
18      for (j in 1:ni) a = a + 1
19      for (j in 1:nd) b = b * 1.00001
20    }
21    return(list(u=u, count=count))
22  },

```

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

my.unif.r

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

```

23 {
24     # C implementation
25     set.seed(seed)
26     out = .C("cpu_runif_wrap",
27             as.integer(n),
28             as.double(ub),
29             as.integer(ni),
30             as.integer(nd),
31             u=double(n),
32             count=integer(n))
33     return(list(u=out$u, count=out$count))
34 },
35 {
36     # GPU implementation
37     out = .C("gpu_runif", as.integer(n), as.double(ub),
38             as.integer(ni), as.integer(nd),
39             as.double(seed),
40             u=double(n), count=integer(n))
41     return(list(u=out$u, count=out$count))
42 })
43 }

```

Running the example

- The files, `comparison.r` and `comparison-analysis.r`, compare the performances of the R, C, and GPU rejection samplers.

```

1 > ls
2 demo inst R README.md src
3 > cd src
4 > make
5 /usr/local/cuda/bin/nvcc -arch=sm_20 -c -I. -I/usr/local/include -I/
  usr/local/cuda/include -I/apps/lib64/R/include -I/usr/local/
  NVIDIA-GPU-Computing-SDK/C/common/inc -Xcompiler -fpic -DRPRINT
  -DNDEBUG cpu_runif.c -o cpu_runif.o
6 ...
7
8 > cd ..
9 > ls
10 demo inst R README.md src
11 > cd demo
12 > ls
13 comparison.R      comparison-analysis.R      segfault.R
14 > R CMD BATCH comparison.R & # do this using screen: it takes a
  couple days unless you modify comparison.R
15 > R CMD BATCH comparison-analysis.R
16 > ls
17 comparison-analysis.R      comparison.csv      comparison.Rout      rejection.
  pdf      segfault.R
18 comparison-analysis.Rout      comparison.R      comparison.tex      Rplots.pdf
  sm.tex

```


Performance: ratios of CPU time to GPU time

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

Outline

Host interface

Device interface

Rejection sampling on the GPU

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

Resources

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

- ▶ Guides:

1. CURAND Guide

- ▶ Code from today:

- ▶ `host_api.cu`
- ▶ `device_api.cu`
- ▶ Dr. Niemi's rejection sampling code

That's all for today.

The CURAND
library

Will Landau

Host interface

Device interface

Rejection sampling
on the GPU

- ▶ Series materials are available at
<http://will-landau.com/gpu>.